

Gautham Hullikunte

<https://www.linkedin.com/in/gautham-hullikunte-1175569a/>

Chennai, TN, India

hgautham93@gmail.com

Summary Professional Software Engineer with **8 years of experience**, specialising in **Data Engineering & Backend Development**. Skilled in building **ETL pipelines, REST APIs and Microservices** using **Java and Python**. Experienced with **cloud platforms (AWS, GCP), distributed systems (Spark, Hadoop), and container orchestration (Docker, Kubernetes)**

Created <https://techmaps.fyi/>

Open source / Side project Portfolio: <https://github.com/batcity> - **700+ commits**

Experience Techmaps.fyi Chennai, India
Founder & Engineer May 2026 - Present

- **Automated Data Discovery:** Engineered an ingestion workflow using Google GenAI SDK (Gemini Flash) and Pydantic for structured outputs, incorporating strict taxonomy rules and rate-limiting logic
- Developed an AI powered platform for exploring authenticated engineering centers across Indian tech hubs including Bengaluru, Chennai, Hyderabad etc
- **Data Validation Guards:** Created human-in-the-loop systems utilizing LLM inference and Google Places API for grounding, domain matching, and fuzzy filtering to ensure location accuracy
- Built a moderation workflow via Streamlit featuring spatial verification and SQLite staging, automating production syncs through GitHub Pull Requests using PyGithub
- Developed a responsive frontend with Next.js 14, TypeScript, and Tailwind CSS, implementing keyboard shortcuts, real-time filtering, and programmatic SEO generation
- Established testing protocols using Vitest and Python unittest, achieving comprehensive coverage across UI components, data parsing, and validation layers

TalentNeuron (Labor Force Analytics Firm) Remote, Chennai, India
Contractor - Senior Software Engineer May 2024 - Mar 2025

- **Automation of Data Processing:** Architected an end-to-end solution reducing a 12-step manual data processing task (millions of records) to a single-click operation using AWS Step Functions, Lambda, AWS Glue & SNS, slashing developer work by 90%
- Developed a **PySpark Data Ingestion Pipeline** for uploading data to multiple databases based on CSV file uploads to designated S3 folders
- **Collaborated with the VP** to create a generic Python-based AWS Lambda function that automates data loading to five different tables based on CSV file uploads to specific S3 folders, this **facilitated data transfer from a recently acquired startup**: <https://www.talentneuron.com/about/newsroom/talentneuron-acquires-hrforecast>
- Implemented a **data alerting system** using Python and Lambda functions, sending SNS alerts to a Teams channel upon arrival of new data in specific S3 folders

Niantic, Inc (AR startup that spun out of Google)

Kansas City, USA

Senior Software Engineer

June 2022 - June 2023

- Built and maintained Data pipelines using **Golang, Python, Argo workflows, Docker and Kubernetes** for Niantic's Visual Positioning System: <https://lightship.dev/products/vps>
- Wrote an **Apache Beam dataflow using Python** to run validation on millions of records
- Worked with the **GCP cloud platform** - specifically with Spanner, BigQuery, GCS, Dataflow and GKE
- **Led cross-team discussions and authored the high-level design for Niantic's Private VPS** (new enterprise feature for secure, access-controlled AR localization), defining APIs, data pipeline updates, and responsibilities across frontend, backend, and data engineering teams
- Implemented a large scale project that improved the quality of 100K+ points of interest on the Niantic AR map (specifically the data pipeline side of things)
- **Worked on Cost estimates** for a large scale data pipeline run that improved the quality of thousands of points of interest on the Niantic AR map

Cerner (Healthcare tech firm acquired by Oracle)

Kansas City, USA

Associate Senior Software Engineer

Sep 2018 – May 2022

Software Engineer

June 2016 – Sep 2018

- Developed and maintained new **batch and incremental data pipelines** that were highly available and scalable using **Java, Apache Crunch, Hadoop and Oozie**
- Improved and maintained an **ETL job** that transfers data from **HDFS to an HBase datastore**
- Developed **REST APIs using an internal Java framework** called beadledom: <https://github.com/cerner/beadledom>
- Developed an automated SLA monitoring system which ran on Jenkins
- **Mentored and helped onboard** Junior Engineers and Interns

Skills

Data Engineering: Spark (PySpark), Hadoop, MySQL, Spanner, BigQuery

Backend Engineering: REST APIs, Microservices, Spring Boot (Basic)

Infrastructure: AWS, GCP, Docker, Kubernetes, Jenkins, Terraform

Projects

- Created a fix for a bug in a REST endpoint to ensure accurate error handling when force-deleting an already deleted model in Elasticsearch: <https://github.com/elastic/elasticsearch/pull/107188>
- Investigated JVM memory pressure from bulk indexing and refactored a core component to simplify object lifetimes: <https://github.com/opensearch-project/OpenSearch/pull/20454>
- Fixed TypeError in reify when folding sparse arrays in Dask Bags (Language: Python)
Github link - <https://github.com/dask/dask/pull/12103>
- Modified Docker compose up command to respect COMPOSE_REMOVE_ORPHANS environment variable, Github link - <https://github.com/docker/compose/issues/11462>

Education

Missouri University of Science and Technology (Rolla, MO, USA)

May 2016

M.S. Computer Science